

# COMP4222 Machine Learning with Structured Data

PageRank

Yangqiu Song

**Slides credits: Jeffrey D. Ullman and Jure Leskovec @Stanford**

# PageRank

- The year 1998 was an eventful year for Web link analysis models. Both the **PageRank** and **HITS** algorithms were reported in that year.
- The **connections** between PageRank and HITS are quite striking.
- Since that eventful year, PageRank has emerged as the dominant link analysis model,
  - due to its query-independence,
  - its ability to combat spamming, and
  - Google's huge business success.

# Ranking web pages

- Web pages are not equally “important”
  - <https://www.facebook.com/> vs <https://www.ust.hk/>
- Inlinks as votes
  - <https://www.facebook.com/> 668.7M backlinks from 609.2K domains
  - <https://www.ust.hk/> 7.1K backlinks from 270 domains
- Are all inlinks equal?
  - Recursive question!




# Simple recursive formulation

- Each link's vote is proportional to the **importance** of its source page
- If page **P** with importance **x** has **n** outlinks, each link gets  **$x/n$**  votes
- Page **P**'s own importance is the sum of the votes on its inlinks

# Simple “flow” model

The web in 1839



$$y = y/2 + a/2$$

$$a = y/2 + m$$

$$m = a/2$$

# Solving the flow equations


- 3 equations, 3 unknowns, no constants
  - No unique solution
  - All solutions equivalent modulo scale factor
- Additional constraint forces uniqueness
  - $y+a+m = 1$
  - $y = 2/5, a = 2/5, m = 1/5$
- Gaussian elimination method works for small examples, but we need a better method for large graphs

# Matrix formulation

- Matrix **M** has one row and one column for each web page
- Suppose page  $j$  has  $n$  outlinks
  - If  $j$  has an outlink  $i$ , then  $M_{ij}=1/n$
  - Else  $M_{ij}=0$
- **M** is a **column stochastic matrix**
  - Columns sum to 1
- Suppose **r** is a vector with one entry per web page
  - $r_j$  is the importance score of page  $j$
  - Call it the **rank vector**
  - $|\mathbf{r}| = 1$

# Example

Suppose page  $j$  links to 3 pages, including  $i$






# Eigenvector Formulation

- The flow equations can be written

$$\mathbf{r} = \mathbf{M}\mathbf{r}$$

- So the rank vector is an eigenvector of the stochastic web matrix
  - In fact, its first or principal eigenvector, with corresponding eigenvalue 1

# Example



$$y = y/2 + a/2$$

$$a = y/2 + m$$

$$m = a/2$$

	y	a	m
y	1/2	1/2	0
a	1/2	0	1
m	0	1/2	0

$$\mathbf{r} = \mathbf{M}\mathbf{r}$$


$$\begin{bmatrix} y \\ a \\ m \end{bmatrix} = \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 1 \\ 0 & 1/2 & 0 \end{bmatrix} \begin{bmatrix} y \\ a \\ m \end{bmatrix}$$

$$y = 2/5, a = 2/5, m = 1/5$$

# Power Iteration method

- Simple iterative scheme (aka **relaxation**)
- Suppose there are  $N$  web pages
- Initialize:  $\mathbf{r}^0 = [1/N, \dots, 1/N]^T$
- Iterate:  $\mathbf{r}^{k+1} = \mathbf{M}\mathbf{r}^k$
- Stop when  $\|\mathbf{r}^{k+1} - \mathbf{r}^k\|_1 < \varepsilon$ 
  - $\|\mathbf{x}\|_1 = \sum_{1 \leq i \leq N} |x_i|$  is the  $L_1$  norm
  - Can use any other vector norm e.g., Euclidean

# Power Iteration Example



	y	a	m
y	$1/2$	$1/2$	0
a	$1/2$	0	1
m	0	$1/2$	0

y		$1/3$	$1/3$	$5/12$	$3/8$		$2/5$
a	=	$1/3$	$1/2$	$1/3$	$11/24$	...	$2/5$
m		$1/3$	$1/6$	$1/4$	$1/6$		$1/5$

# Random Walk Interpretation

- Imagine a **random web surfer**
  - At any time  $t$ , surfer is on some page  $P$
  - At time  $t+1$ , the surfer follows an outlink from  $P$  uniformly at random
  - Ends up on some page  $Q$  linked from  $P$
  - Process repeats indefinitely
- Let  $\mathbf{p}(t)$  be a vector whose  $i^{\text{th}}$  component is the probability that the surfer is at page  $i$  at time  $t$ 
  - $\mathbf{p}(t)$  is a probability distribution on pages

# The stationary distribution

- Where is the surfer at time  $t+1$ ?
  - Follows a link uniformly at random
  - $\mathbf{p}(t+1) = \mathbf{M}\mathbf{p}(t)$
- Suppose the random walk reaches a state such that  $\mathbf{p}(t+1) = \mathbf{M}\mathbf{p}(t) = \mathbf{p}(t)$ 
  - Then  $\mathbf{p}(t)$  is called a stationary distribution for the random walk
- Our rank vector  $\mathbf{r}$  satisfies  $\mathbf{r} = \mathbf{M}\mathbf{r}$ 
  - So it is a stationary distribution for the random surfer

# Existence and Uniqueness

A central result from the theory of random walks (aka Markov processes):


For graphs that satisfy certain conditions, the stationary distribution is unique and eventually will be reached no matter what the initial probability distribution at time  $t = 0$ .

Problems?




# Spider traps

- A group of pages is a **spider trap** if there are no links from within the group to outside the group
  - Random surfer gets trapped
- Spider traps violate the conditions needed for the random walk theorem




# Microsoft becomes a spider trap



	y	a	m
y	1/2	1/2	0
a	1/2	0	0
m	0	1/2	1


y	=	1	1	3/4	5/8		0
a		1	1/2	1/2	3/8	...	0
m		1	3/2	7/4	2		3

# Random Teleports



- The Google solution for spider traps
- At each time step, the random surfer has two options:
  - With probability  $\beta$ , follow a link at random
  - With probability  $1-\beta$ , jump to some page uniformly at random
  - Common values for  $\beta$  are in the range 0.8 to 0.9
- Surfer will teleport out of spider trap within a few time steps

# Random teleports ( $\beta = 0.8$ )




$$\begin{array}{c} y \\ a \\ m \end{array} \begin{array}{c} y \\ 1/2 \\ 1/2 \\ 0 \end{array} \quad 0.8 * \begin{array}{c} y \\ 1/2 \\ 1/2 \\ 0 \end{array} + 0.2 * \begin{array}{c} y \\ 1/3 \\ 1/3 \\ 1/3 \end{array}$$

$$0.8 \begin{array}{ccc} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 0 \\ 0 & 1/2 & 1 \end{array} + 0.2 \begin{array}{ccc} 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \end{array}$$

$$\begin{array}{c} y \\ a \\ m \end{array} \begin{array}{ccc} 7/15 & 7/15 & 1/15 \\ 7/15 & 1/15 & 1/15 \\ 1/15 & 7/15 & 13/15 \end{array}$$

# Random teleports ( $\beta = 0.8$ )



$$0.8 \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 0 \\ 0 & 1/2 & 1 \end{bmatrix} + 0.2 \begin{bmatrix} 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \end{bmatrix}$$

$$\begin{matrix} y \\ a \\ m \end{matrix} \begin{bmatrix} 7/15 & 7/15 & 1/15 \\ 7/15 & 1/15 & 1/15 \\ 1/15 & 7/15 & 13/15 \end{bmatrix}$$


$$\begin{matrix} y \\ a \\ m \end{matrix} = \begin{matrix} 1 & 1.00 & 0.84 & 0.776 & & 7/11 \\ 1 & 0.60 & 0.60 & 0.536 & \dots & 5/11 \\ 1 & 1.40 & 1.56 & 1.688 & & 21/11 \end{matrix}$$

# Page Rank


- Construct the  $N \times N$  matrix  $\mathbf{A}$  as follows
  - $A_{ij} = \beta M_{ij} + (1-\beta)/N$
- Verify that  $\mathbf{A}$  is a stochastic matrix
- The **page rank vector**  $\mathbf{r}$  is the principal eigenvector of this matrix
  - satisfying  $\mathbf{r} = \mathbf{A}\mathbf{r}$
- Equivalently,  $\mathbf{r}$  is the stationary distribution of the random walk with teleports

# Dead ends

- Pages with no outlinks are “dead ends” for the random surfer
  - Nowhere to go on next step



# Microsoft becomes a dead end



$$0.8 \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 0 \\ 0 & 1/2 & 0 \end{bmatrix} + 0.2 \begin{bmatrix} 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \end{bmatrix}$$

$$\begin{matrix} y \\ a \\ m \end{matrix} \begin{bmatrix} 7/15 & 7/15 & 1/15 \\ 7/15 & 1/15 & 1/15 \\ 1/15 & 7/15 & 1/15 \end{bmatrix}$$




Non-stochastic!  
(Sum of column is not 1)

$$\begin{matrix} y \\ a \\ m \end{matrix} = \begin{matrix} 1 & 1 & 0.787 & 0.648 & & 0 \\ 1 & 0.6 & 0.547 & 0.430 & \dots & 0 \\ 1 & 0.6 & 0.387 & 0.333 & & 0 \end{matrix}$$





# Solution to Dead Ends

- **Teleports:** Follow random teleport links with total probability **1.0** from dead-ends
  - Adjust matrix accordingly



	y	a	m
y	$\frac{1}{2}$	$\frac{1}{2}$	0
a	$\frac{1}{2}$	0	0
m	0	$\frac{1}{2}$	0



	y	a	m
y	$\frac{1}{2}$	$\frac{1}{2}$	$\frac{1}{3}$
a	$\frac{1}{2}$	0	$\frac{1}{3}$
m	0	$\frac{1}{2}$	$\frac{1}{3}$

# Why Teleports Solve the Problem?

- Why are dead-ends and spider traps a problem and why do teleports solve the problem?
- **Spider-traps** are not a problem, but with traps PageRank scores are **not** what we want
  - **Solution:** Never get stuck in a spider trap by teleporting out of it in a finite number of steps
- **Dead-ends** are a problem
  - The matrix is not column stochastic so our initial assumptions are not met
  - **Solution:** Make matrix column stochastic by always teleporting when there is nowhere else to go

# Computational Issues

# Computing PageRank

- Key step is matrix-vector multiplication
  - $\mathbf{r}^{\text{new}} = \mathbf{A}\mathbf{r}^{\text{old}}$
- Easy if we have enough main memory to hold  $\mathbf{A}$ ,  $\mathbf{r}^{\text{old}}$ ,  $\mathbf{r}^{\text{new}}$
- Say  $N = 1$  billion pages
  - We need 4 bytes for each entry (say)
  - 2 billion entries for vectors, approx 8GB
  - Matrix  $\mathbf{A}$  has  $N^2$  entries
    - $10^{18}$  is a large number!

# Rearranging the Equation

$\mathbf{r} = \mathbf{Ar}$ , where

$$A_{ij} = \beta M_{ij} + (1-\beta)/N$$

$$r_i = \sum_{1 \leq j \leq N} A_{ij} r_j$$

$$\begin{aligned} r_i &= \sum_{1 \leq j \leq N} [\beta M_{ij} + (1-\beta)/N] r_j \\ &= \beta \sum_{1 \leq j \leq N} M_{ij} r_j + (1-\beta)/N \sum_{1 \leq j \leq N} r_j \\ &= \beta \sum_{1 \leq j \leq N} M_{ij} r_j + (1-\beta)/N, \text{ since } |\mathbf{r}| = 1 \end{aligned}$$

$$\mathbf{r} = \beta \mathbf{M}\mathbf{r} + [(1-\beta)/N]_N$$

where  $[x]_N$  is an  $N$ -vector with all entries  $x$

$$0.8 \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 0 \\ 0 & 1/2 & 1 \end{bmatrix} + 0.2 \begin{bmatrix} 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \end{bmatrix}$$

y	7/15	7/15	1/15
a	7/15	1/15	1/15
m	1/15	7/15	13/15

# Sparse Matrix Formulation

- We can rearrange the page rank equation:
  - $\mathbf{r} = \beta \mathbf{M} \mathbf{r} + [(1-\beta)/N]_N$
  - $[(1-\beta)/N]_N$  is an N-vector with all entries  $(1-\beta)/N$
- $\mathbf{M}$  is a sparse matrix!
  - 10 links per node, approx  $10N$  entries
- So in each iteration, we need to:
  - Compute  $\mathbf{r}^{\text{new}} = \beta \mathbf{M} \mathbf{r}^{\text{old}}$
  - Add a constant value  $(1-\beta)/N$  to each entry in  $\mathbf{r}^{\text{new}}$

# Sparse Matrix Encoding

- Encode sparse matrix using only nonzero entries
  - Space proportional roughly to number of links
  - Say  $10N$ , or  $4 \times 10^1$  billion = 40GB

source node	degree	destination nodes
0	3	1, 5, 7
1	5	17, 64, 113, 117, 245
2	2	13, 23

# Basic Algorithm

- Initialize:  $\mathbf{r}^{\text{old}} = [1/N]_N$
- Iterate:
  - **Update**: Perform a sequential scan of  $\mathbf{M}$  and  $\mathbf{r}^{\text{old}}$  to update  $\mathbf{r}^{\text{new}}$
  - Every few iterations, compute  $|\mathbf{r}^{\text{new}} - \mathbf{r}^{\text{old}}|$  and stop if it is below threshold



# Summary

- We introduced
  - Network analysis, centrality
  - PageRank, which powers Google
- **Important to note:** Hyperlink based ranking is not the only algorithm used in search engines. In fact, it is combined with many **content based factors** to produce the final ranking presented to the user.