

COMP4222 Machine Learning with Structured Data


Graph Neural Networks 1

Instructor: Yangqiu Song

Slides credits: Jure Leskovec, William L. Hamilton, Rex Ying, Rok Sosic

Embedding Nodes


- Goal is to encode nodes so that **similarity in the embedding space (e.g., dot product)** approximates **similarity in the original network**.



Embedding Nodes

Goal: $\text{similarity}(u, v) \approx \mathbf{z}_v^\top \mathbf{z}_u$

Need to define!



Two Key Components

- **Encoder** maps each node to a low-dimensional vector.

$$\text{ENC}(v) = \mathbf{z}_v$$

node in the input graph

d-dimensional embedding

- **Similarity function** specifies how relationships in vector space map to relationships in the original network.

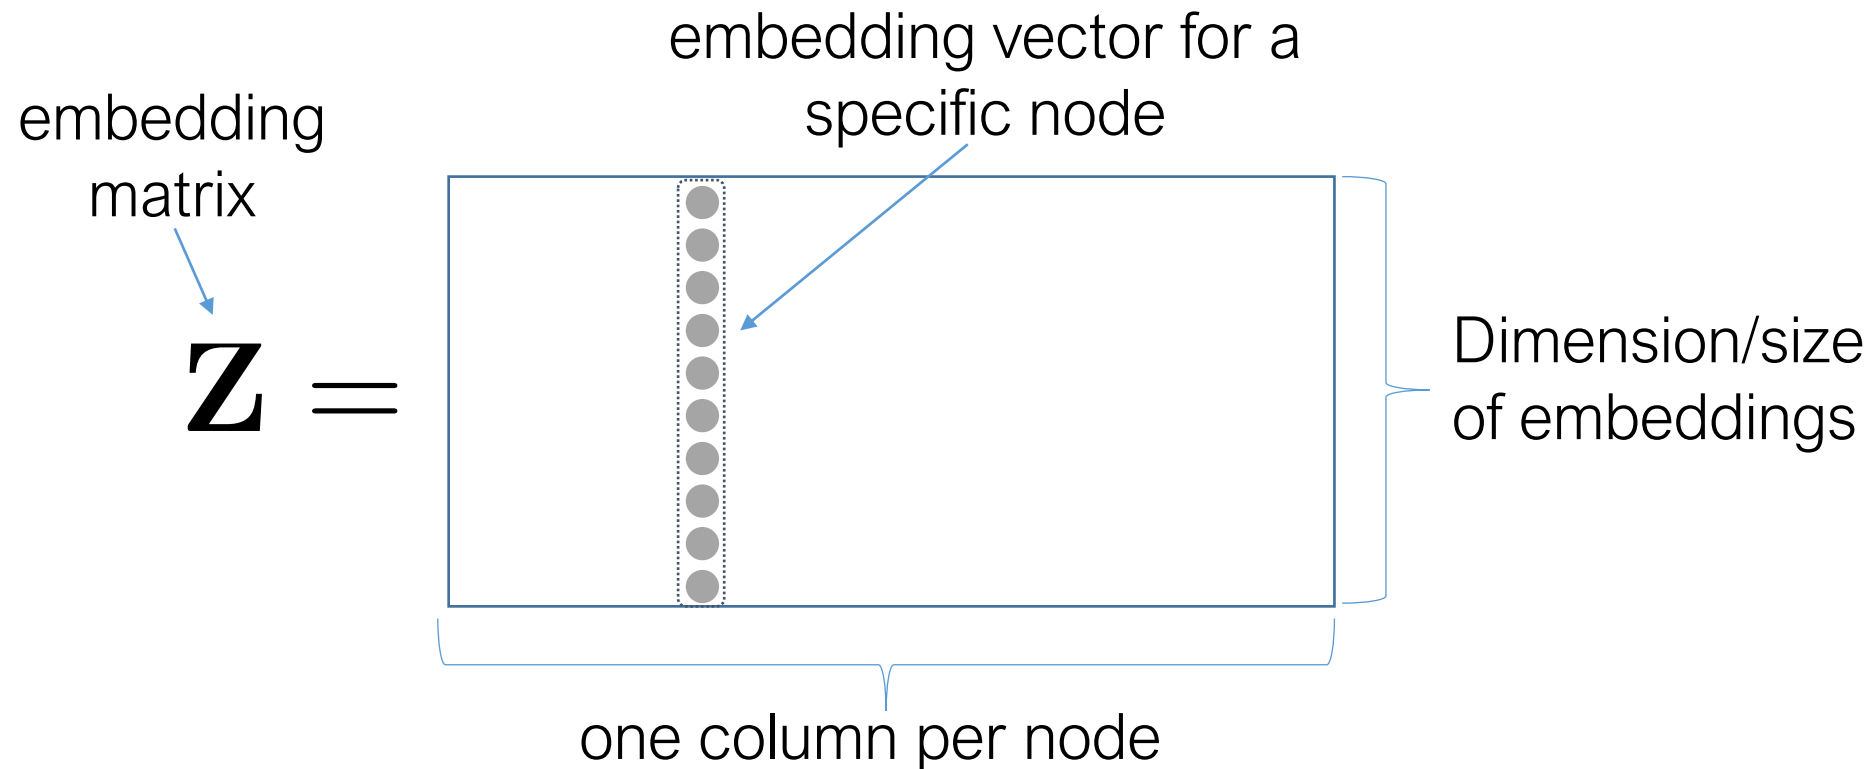
$$\text{similarity}(u, v) \approx \mathbf{z}_v^\top \mathbf{z}_u$$

Similarity of u and v in the original network

dot product between node embeddings

Shallow Encoders

- So far we have focused on “**shallow**” encoders, i.e. embedding lookups:



Shallow Encoders

○ Limitations of shallow embedding methods:

- **$O(|V|)$ parameters are needed:**
 - No sharing of parameters between nodes
 - Every node has its own unique embedding
- **Inherently “transductive”:**
 - Cannot generate embeddings for nodes that are not seen during training
- **Do not incorporate node features:**
 - Many graphs have features that we can and should leverage

From “Shallow” to “Deep”


- We will now discuss “deeper” methods based on **graph neural networks**.

$$\text{ENC}(v) =$$

- complex function that depends on graph structure
- multiple layers of non-linear transformations based on graph structure

- In general, **all of these more complex encoders can be combined with the similarity functions from the previous section.**

Deep Graph Encoders



Tasks on Networks

○ Node classification

- Predict a type of a given node

○ Link prediction

- Predict whether two nodes are linked

○ Community detection

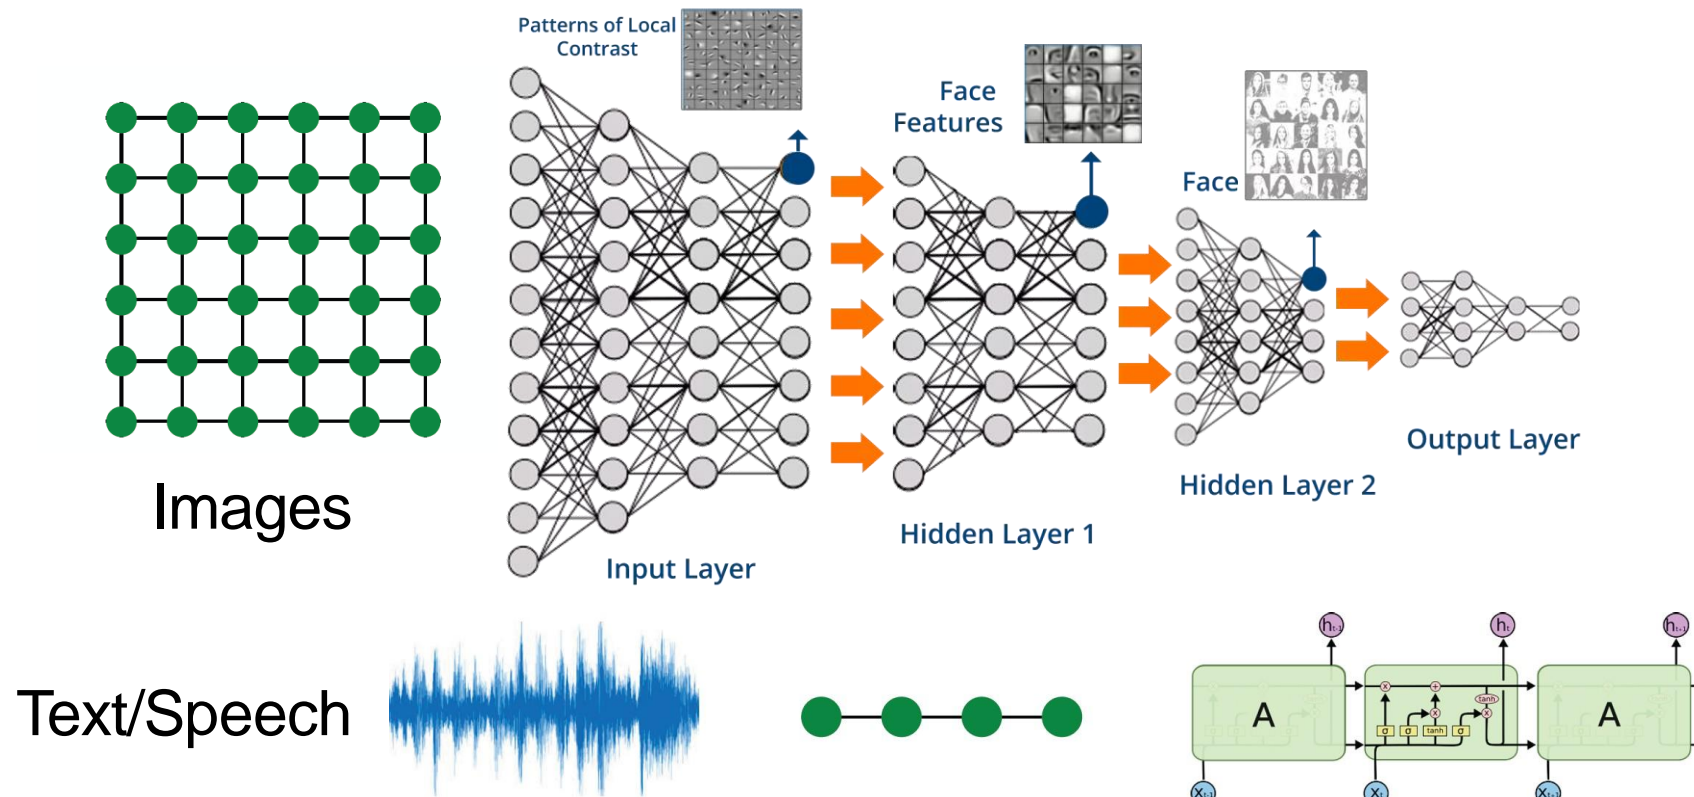
- Identify densely linked clusters of nodes

○ Network similarity

- How similar are two (sub)networks


Modern Machine Learning

- Modern deep learning toolbox is designed for simple sequences & grids




Networks are Far More Complex!

- Arbitrary size and complex topological structure (i.e., no spatial locality like grids)
 - No fixed node ordering or reference point
 - Often dynamic and have multimodal features




Real-World Graphs

- **But our graphs look like this:**




or this:




- There is no fixed notion of locality or sliding window on the graph
- Graph is permutation invariant

From Images to Graphs

- Single CNN layer with 3x3 filter



Image

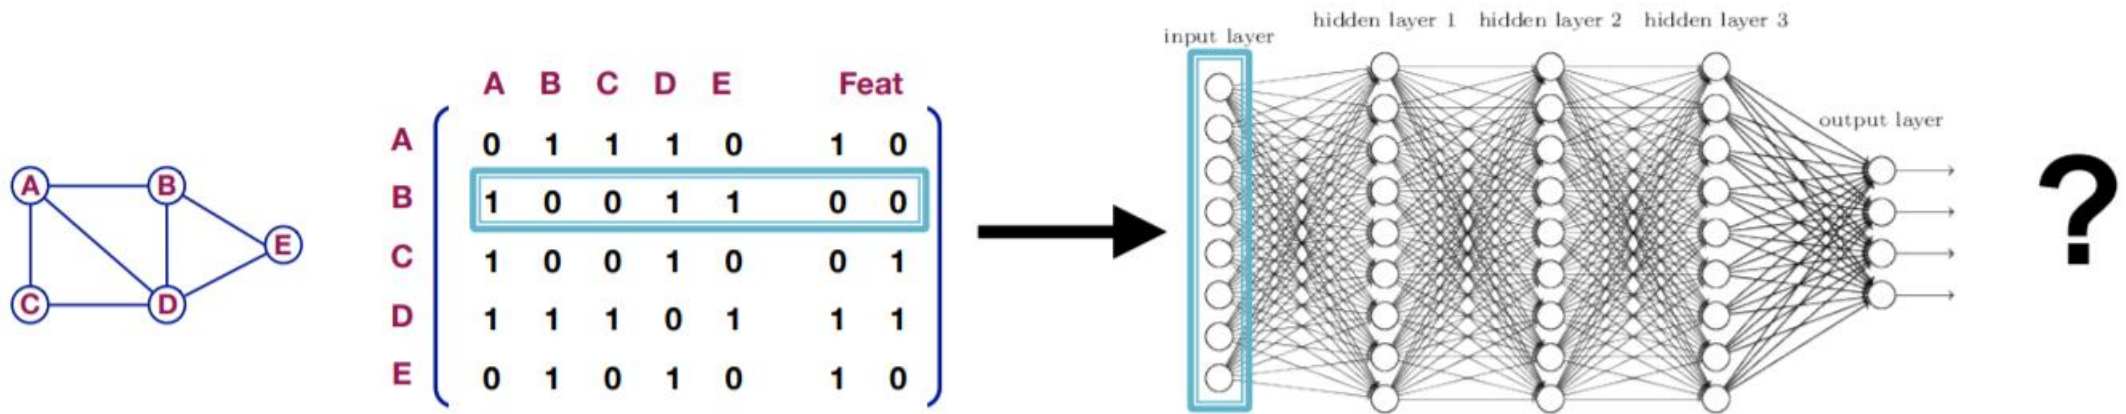


Graph

- Transform information at the neighbors and combine it:
 - Transform “messages” h_i from neighbors: $W_i h_i$
 - Add them up: $\sum_i W_i h_i$

A Naïve Approach

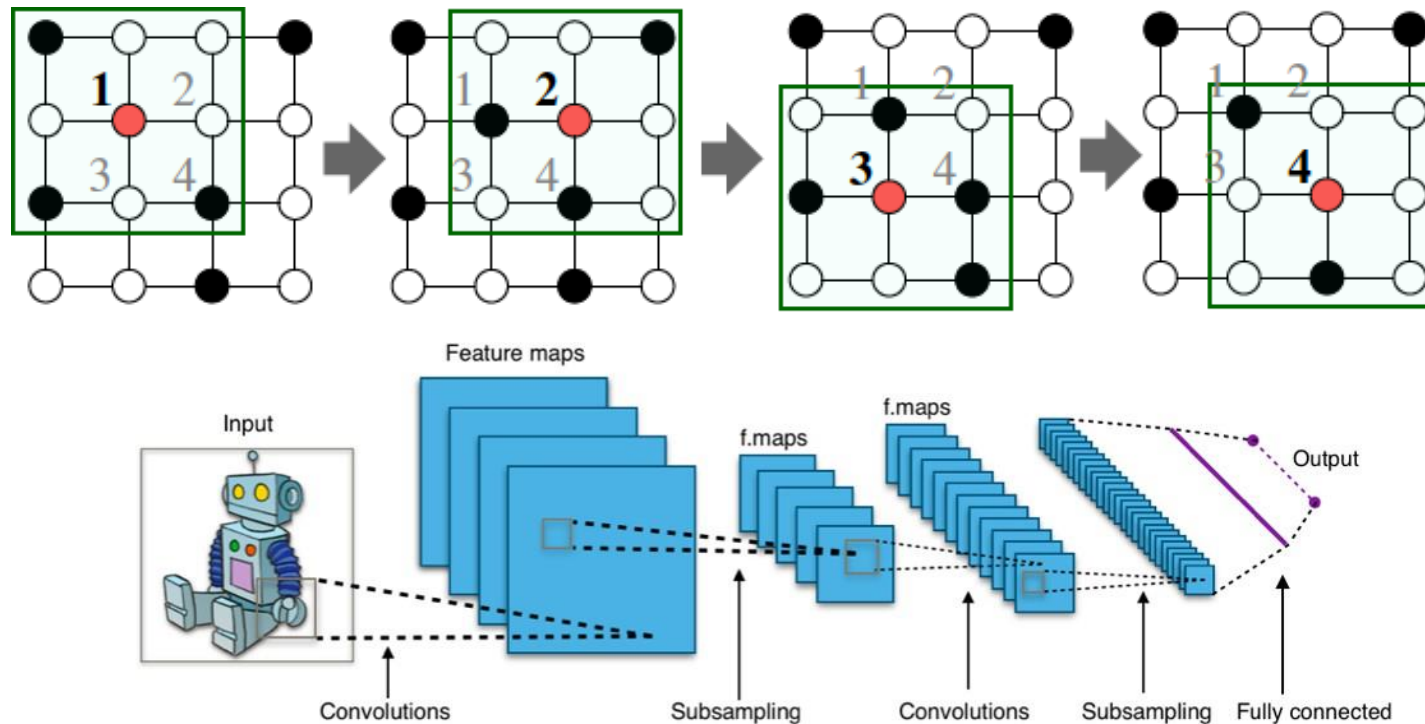
- Join adjacency matrix and features
- Feed them into a deep neural net:



- Issues with this idea:
 - $O(N)$ parameters
 - Not applicable to graphs of different sizes
 - Not invariant to node ordering

General Idea: Graph Convolution

- **CNN on an image:**



Goal is to generalize convolutions beyond simple lattices
Leverage node features/attributes (e.g., text, images)


Setup

- Assume we have a graph G :
 - V is the vertex set.
 - A is the adjacency matrix (assume binary).
 - v is a node in V ; $N(v)$ is the set of neighbors of v
 - $X \in \mathbb{R}^{m \times |V|}$ is a matrix of node features.
 - Attributes, text, image data, E.g.,
 - Social network: user profile, user image
 - Biological network: gene expression profiles, functional information
 - Node degrees, clustering coefficients, etc.
 - Indicator vectors (i.e., one-hot encoding of each node)
 - Retrieve from an embedding matrix which is learnable


Permutation Invariance

- **Graph does not have a canonical order of the nodes!**
- We can have many different order plans.

Order plan 1




Node features X_1




Adjacency matrix A_1

	A	B	C	D	E	F
A						
B						
C						
D						
E						
F						

Order plan 2



Node features X_2



Adjacency matrix A_2

	A	B	C	D	E	F
A						
B						
C						
D						
E						
F						


Graph and node representations should be the same for **Order plan 1** and **Order plan 2**

Permutation Invariance

- What does it mean by “graph representation is same for two order plans”?
- Consider we learn a function f that maps a graph $G = (A, X)$ to a vector \mathbb{R}^d then


$$f(A_1, X_1) = f(A_2, X_2)$$

Order plan 1: A_1, X_1



For two order plans,
output of f should
be the same!

Order plan 2: A_2, X_2



Permutation Invariance

- What does it mean by “graph representation is same for two order plans”?

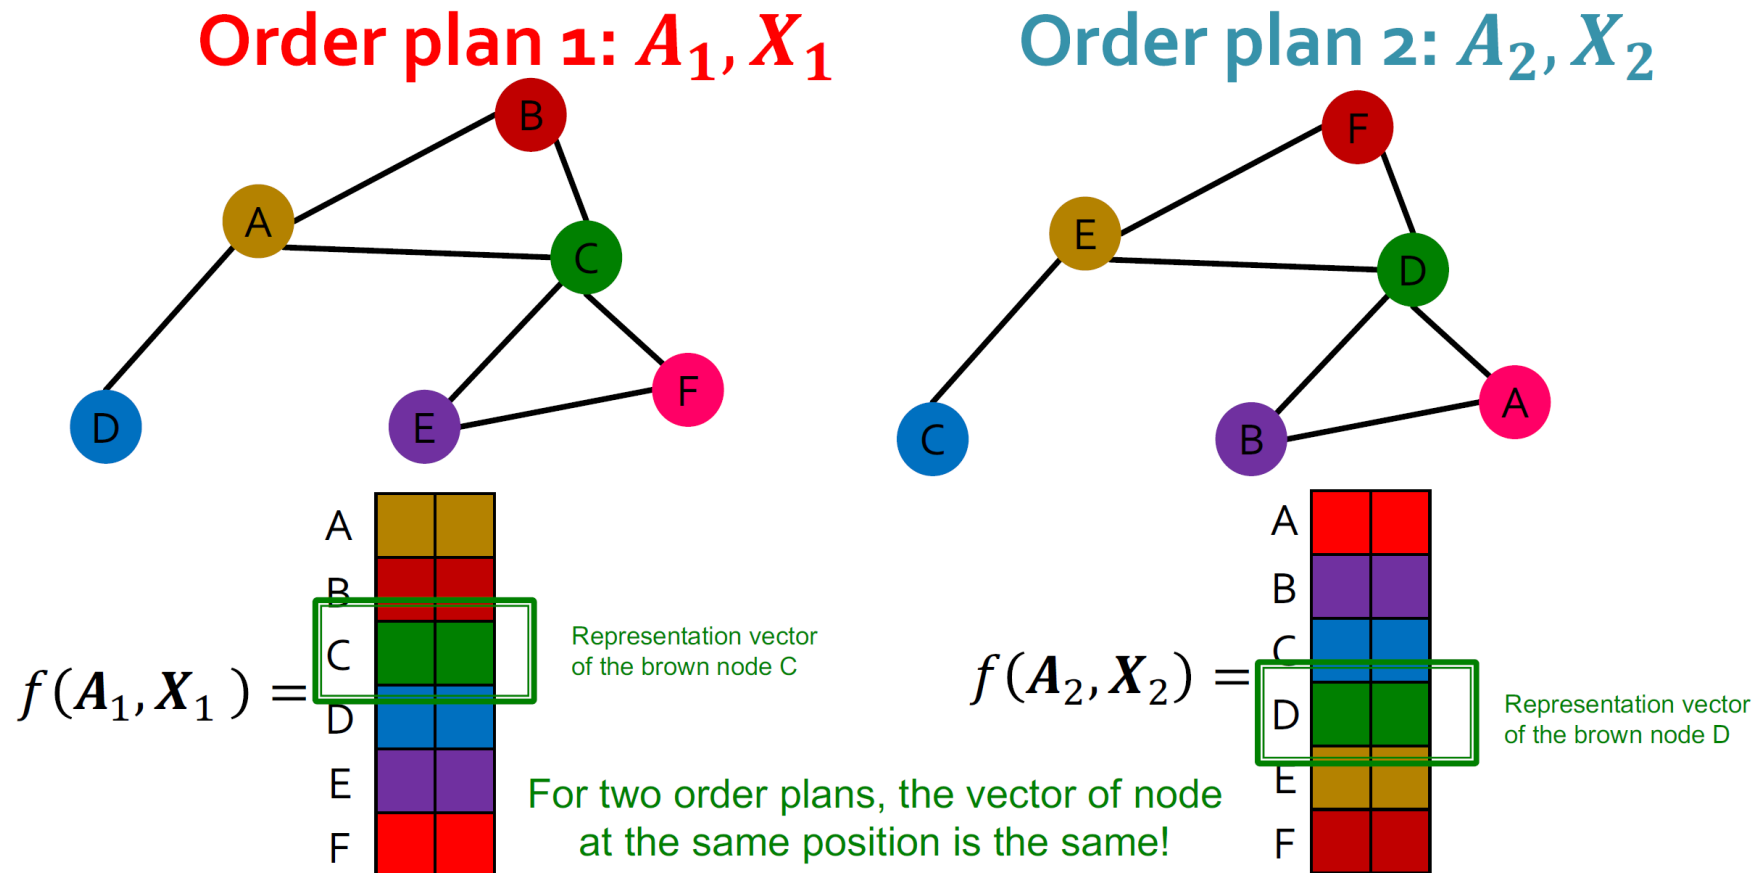
- Consider we learn a function f that maps a graph $G = (A, X)$ to a vector \mathbb{R}^d then

$$f(A_1, X_1) = f(A_2, X_2)$$

- Then if $f(A_i, X_i) = f(A_j, X_j)$ for any order plan i and j , we formally say f is a **permutation invariant function**.
 - For a graph with m nodes, there are $m!$ different order plans.

Permutation Equivariance

- **Similarly for node representation:** We learn a function f that maps nodes of G to a matrix $\mathbb{R}^{m \times d}$



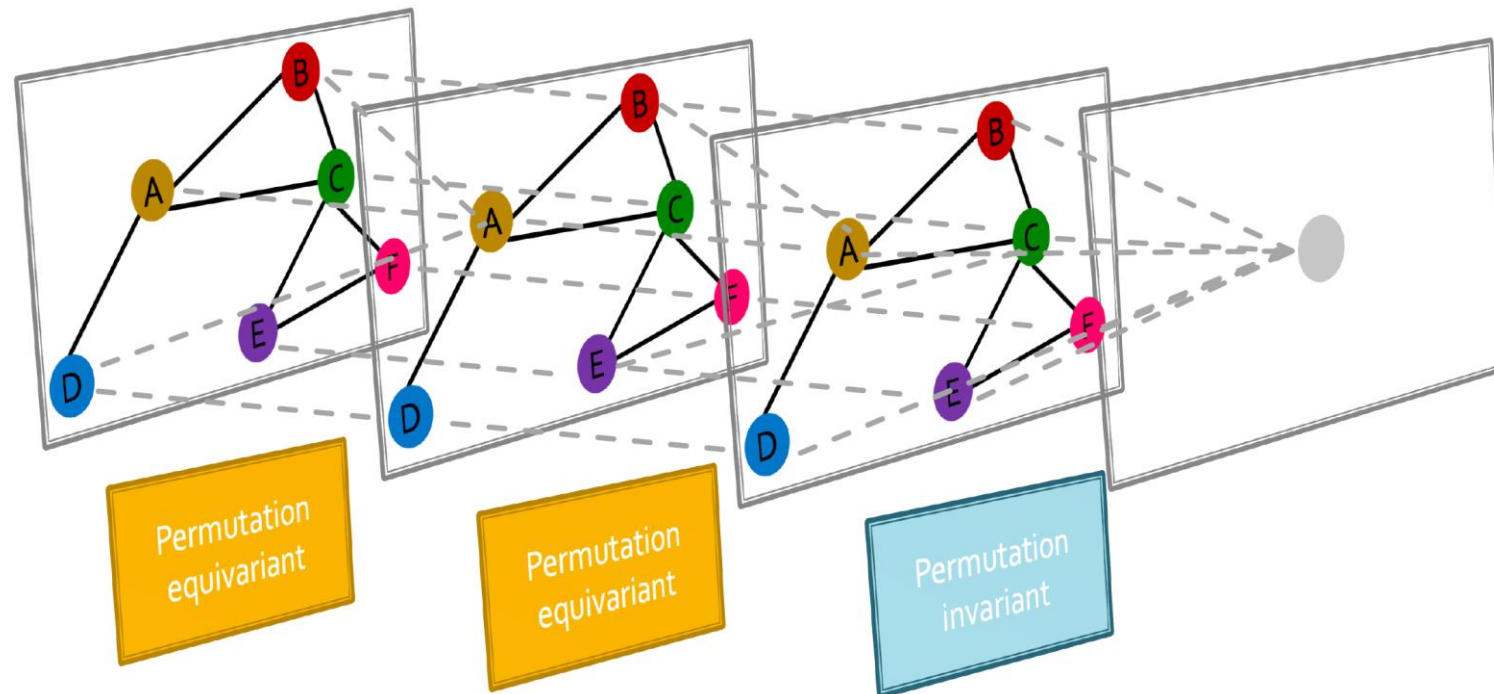
Permutation Equivariance

- **For node representation**

- Consider we learn a function f that maps a graph $G = (\mathbf{A}, \mathbf{X})$ to a matrix $\mathbb{R}^{m \times d}$
 - graph has m nodes, each row is the embedding of a node.
 - For two order plans, the vector of node at the same position is the same!
- Similarly, if this property holds for any pair of order plan i and j , we say f is a **permutation equivariant function**.


Graph Neural Network Overview

- Graph neural networks consist of multiple permutation equivariant / invariant functions



Graph Neural Network Overview

- Are other neural network architectures, e.g., MLPs, permutation invariant / equivariant?
 - No.



Switching the order of the input leads to different outputs!

Outline for this Section

- We will now discuss “deeper” methods based on **graph neural networks**.
 1. The Basics
 2. GraphSAGE


The Basics: Graph Neural Networks

Based on material from:


- Hamilton et al. 2017. [Representation Learning on Graphs: Methods and Applications](#). *IEEE Data Engineering Bulletin on Graph Systems*.
- Scarselli et al. 2005. [The Graph Neural Network Model](#). *IEEE Transactions on Neural Networks*.

Graph Convolutional Networks

- **Idea**: Node's neighborhood defines a computation graph



Determine node
computation graph

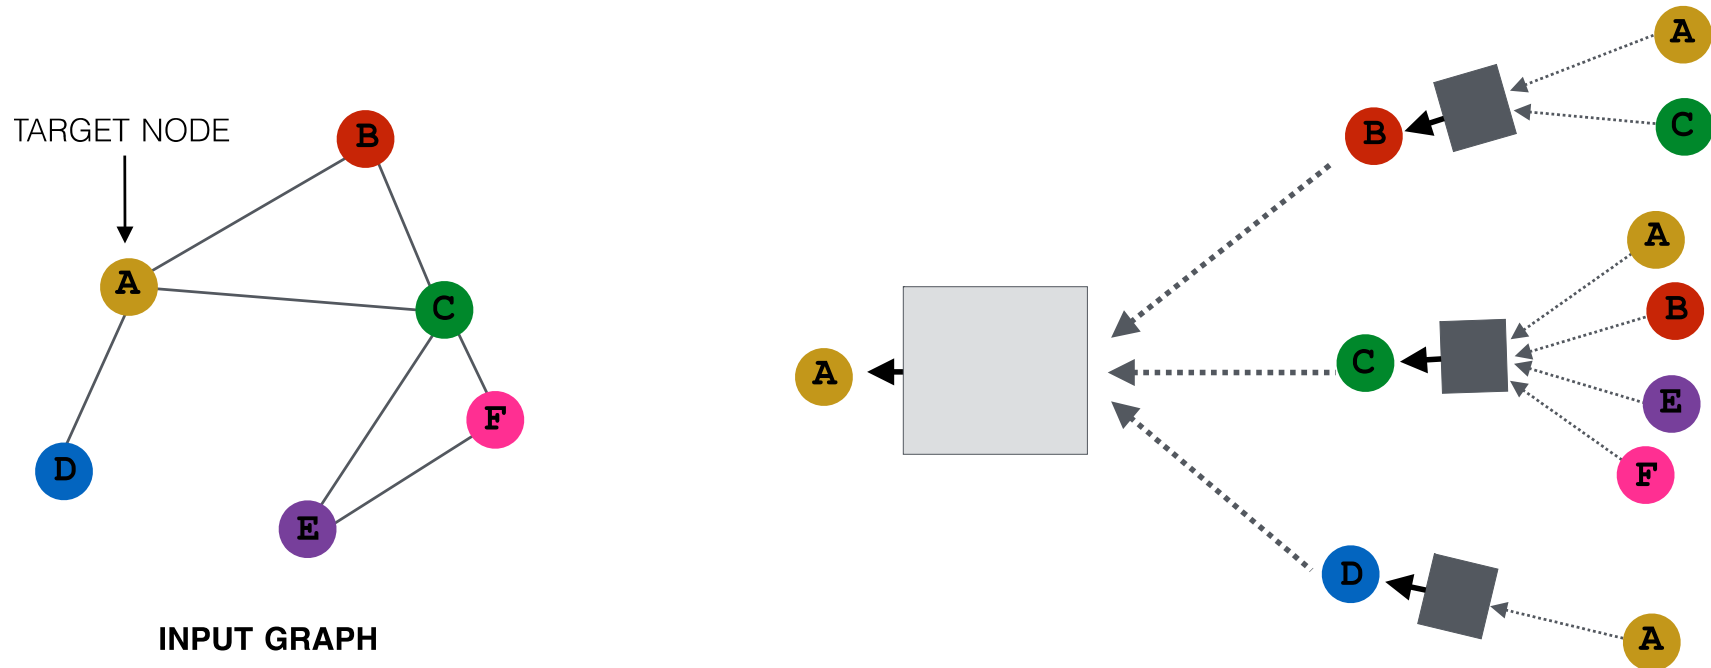


Propagate and
transform information

Learn how to propagate information across the graph to compute node features

Neighborhood Aggregation

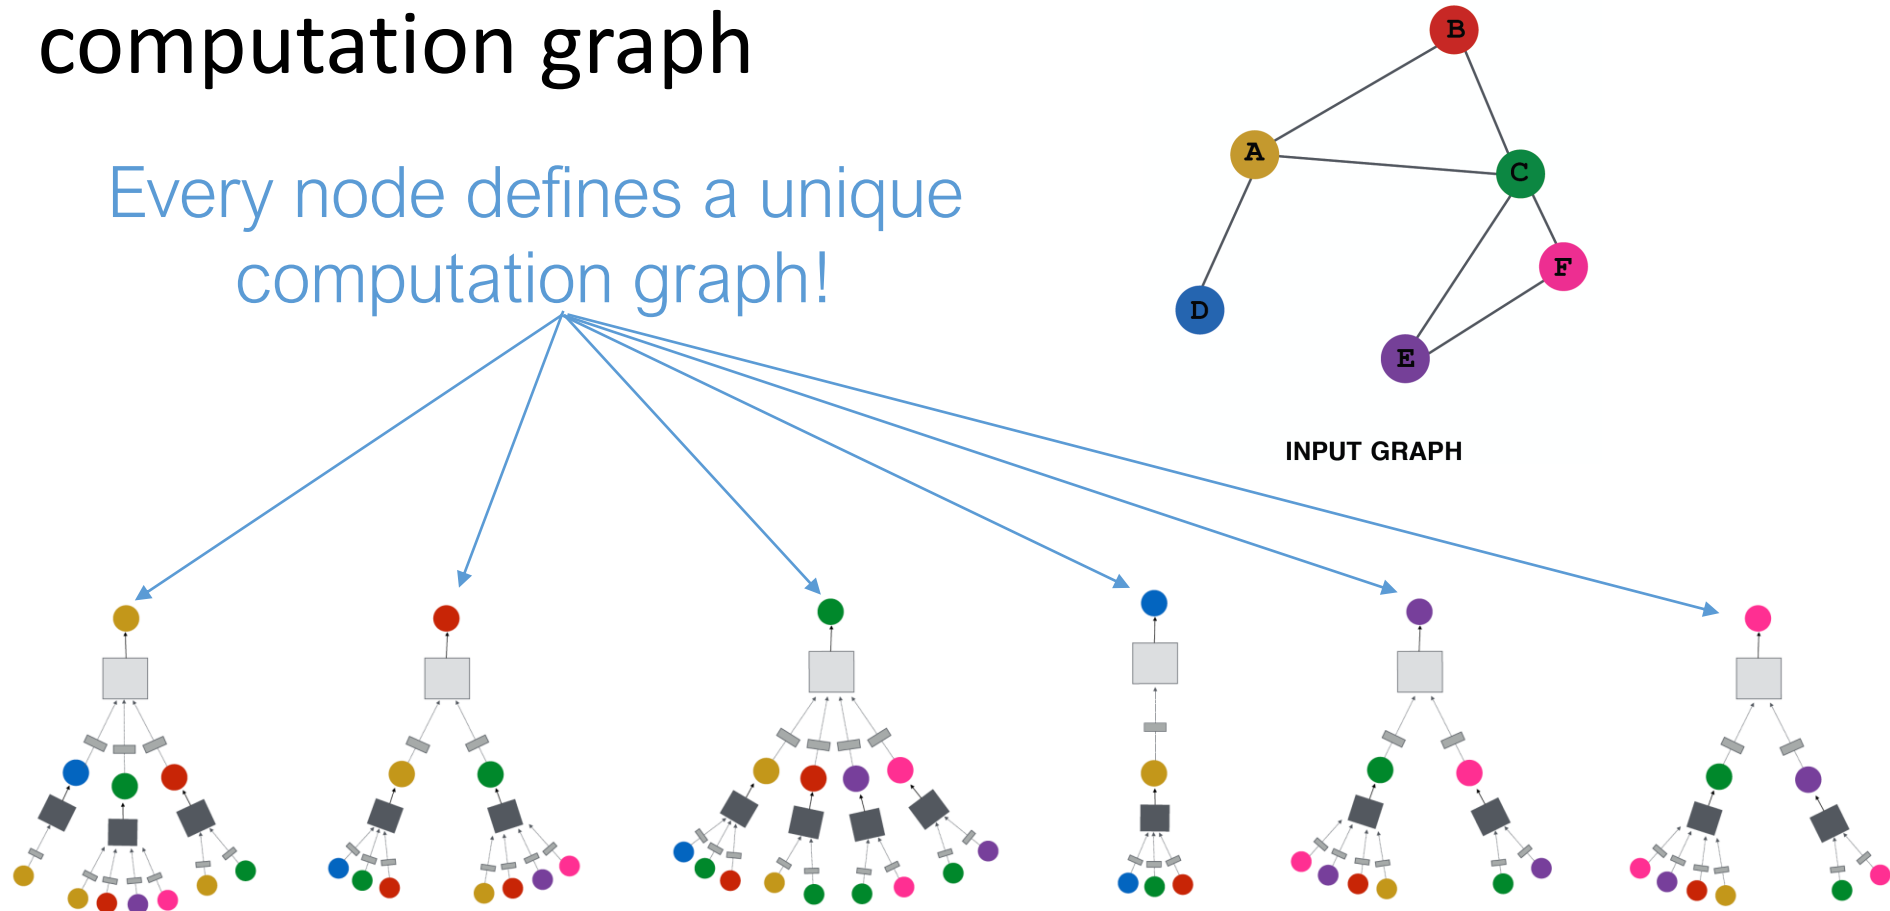
- **Key idea:** Generate node embeddings based on local neighborhoods.
- **Intuition:** Nodes aggregate information from their neighbors using neural networks



Neighborhood Aggregation


- **Intuition:** Network neighborhood defines a computation graph

Every node defines a unique computation graph!




Neighborhood Aggregation

- Nodes have embeddings at each layer.
- Model can be arbitrary depth.
- “Layer-0” embedding of node u is its input feature, i.e. x_u .
- Layer- K embedding gets information from nodes that are K hops away

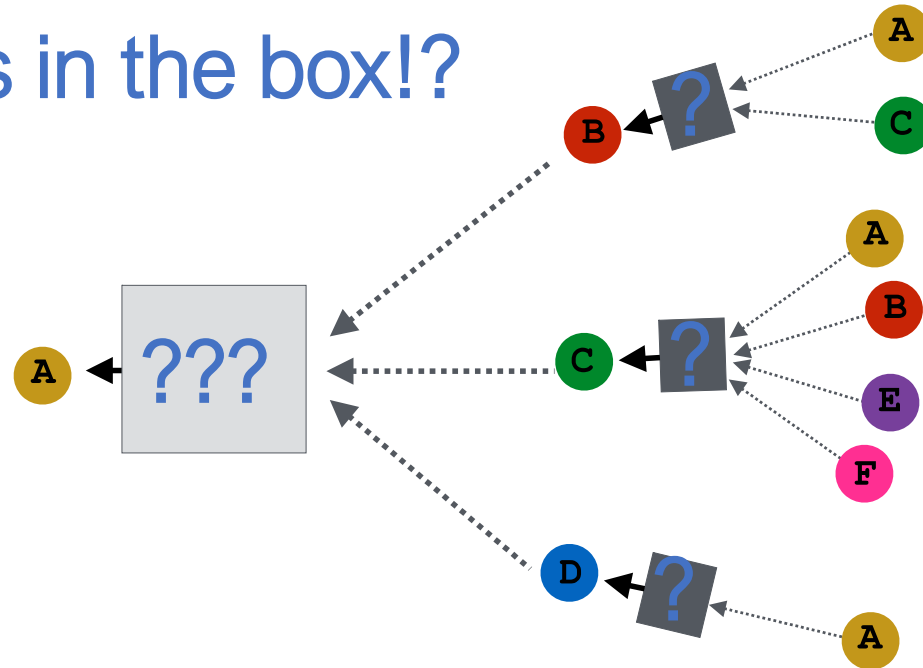


Neighborhood Aggregation

- Key distinctions are in how different approaches aggregate information across the layers.




what's in the box!?



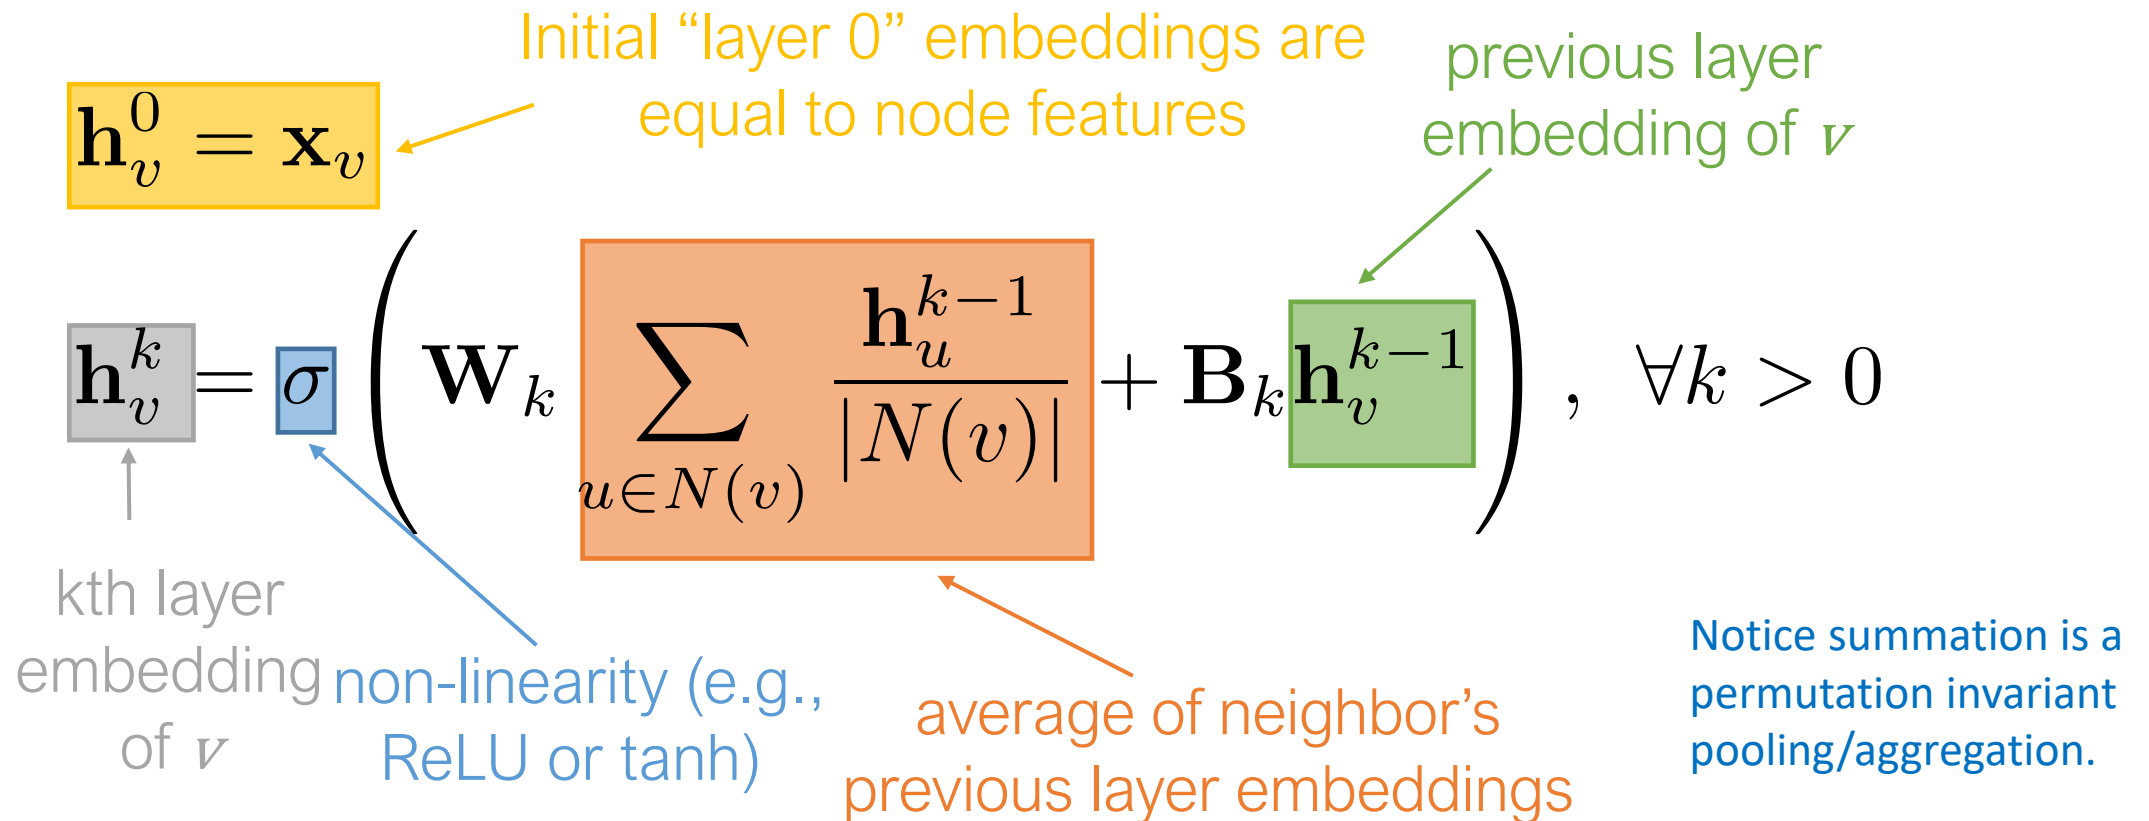
Neighborhood Aggregation

- **Basic approach:** Average neighbor information and apply a neural network.





The Math

- **Basic approach:** Average neighbor messages and apply a neural network.



Equivariant Property

- Message passing and neighbor aggregation in graph convolution networks is permutation equivariant



Node feature X_1

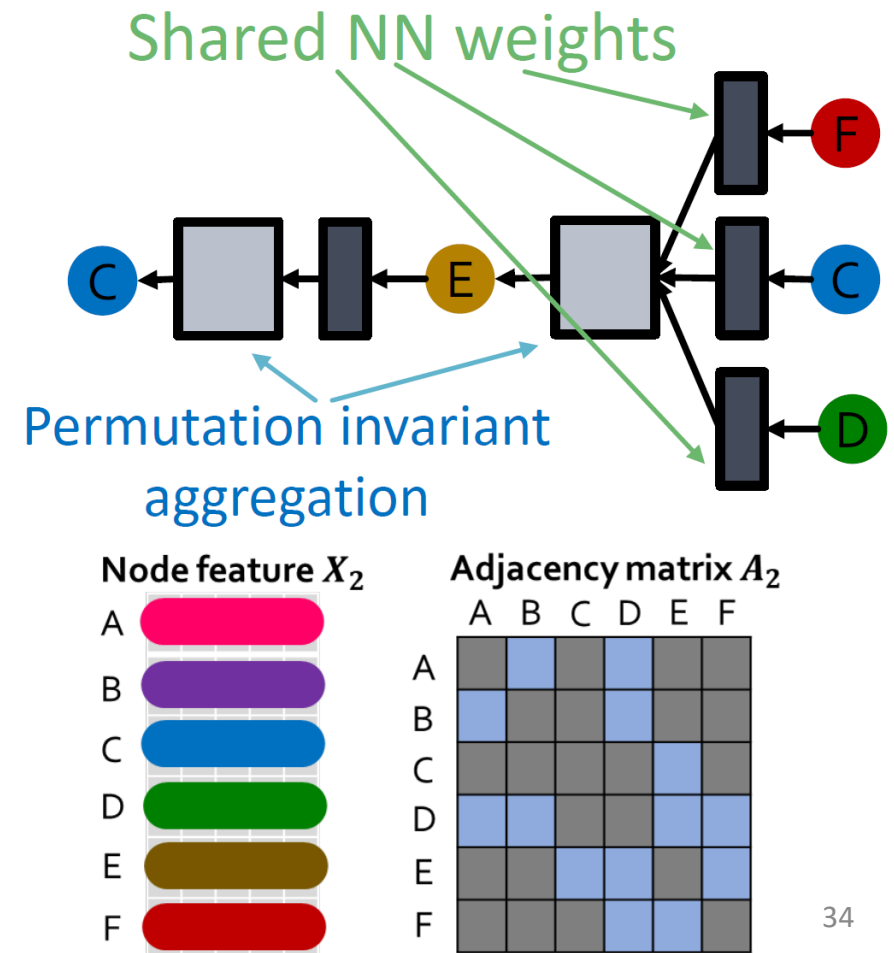
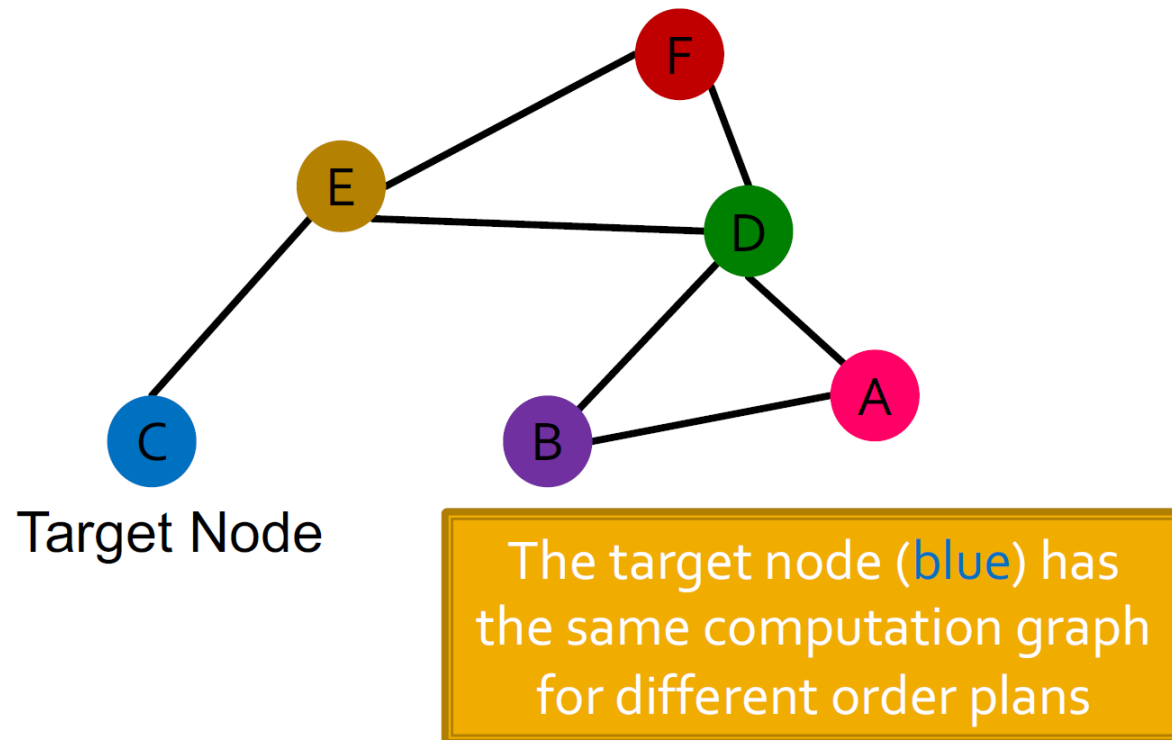
A	
B	
C	
D	
E	
F	

Adjacency matrix A_1

	A	B	C	D	E	F
A						
B						
C						
D						
E						
F						

Equivariant Property

- Message passing and neighbor aggregation in graph convolution networks is permutation equivariant




Matrix Formulation

- **Many aggregations can be performed efficiently by (sparse) matrix operations**

- Let $H^k = [h_1^k, \dots, h_V^k]^T$
- Then $\sum_{u \in N_v} h_u^k = A_{v,:} H^k$
 - N_v is the set of neighbours of v
- Let D be diagonal matrix where
 - $D_{v,v} = \text{Degree}(v) = |N_v|$
 - The inverse of D is also a diagonal matrix
 - $D_{v,v}^{-1} = 1/|N_v|$
- Therefore

$$\sum_{u \in N_v} \frac{h_u^k}{|N_v|} \quad \longrightarrow \quad H^{k+1} = D^{-1} A H^k$$

Matrix of hidden embeddings $H^{(k-1)}$




$h_i^{(k-1)}$

Matrix Formulation


- Re-writing update function in matrix form

$$H^{k+1} = \sigma(D^{-1}AH^k W_k^T + H^k B_k^T)$$

- Red: neighbourhood aggregation
 - Blue: self transformation
-
- In practice, this implies that efficient sparse matrix multiplication can be used (A is sparse)
-
- **Note:** not all GNNs can be expressed in matrix form, when aggregation function is complex




Matrix of hidden embeddings $H^{(k-1)}$



Training the Model

- How do we train the model to generate “high-quality” embeddings?



Need to define a loss function on the embeddings, $\mathcal{L}(z_A)$!

Training the Model

\mathbf{W}_k : weight matrix for neighborhood aggregation
 \mathbf{B}_k : weight matrix for transforming hidden vector of self

trainable matrices
(i.e., what we learn)

$$\mathbf{h}_v^0 = \mathbf{x}_v$$
$$\mathbf{h}_v^k = \sigma \left(\boxed{\mathbf{W}_k} \sum_{u \in N(v)} \frac{\mathbf{h}_u^{k-1}}{|N(v)|} + \boxed{\mathbf{B}_k} \mathbf{h}_v^{k-1} \right), \quad \forall k \in \{1, \dots, K\}$$

$$\boxed{\mathbf{z}_v = \mathbf{h}_v^K}$$


- After K-layers of neighborhood aggregation, we get output embeddings for each node.
- **We can feed these embeddings into any loss function** and run stochastic gradient descent to train the aggregation parameters.

Training the Model


- Train in an **unsupervised manner** using only the graph structure.
 - Use only the graph structure
 - “Similar” nodes have similar embeddings
- Unsupervised loss function can be anything from the last section, e.g., based on
 - Random walks (node2vec, DeepWalk)
 - Adjacent neighbor similarity
 - Node proximity in the graph (higher order relation similarity)

Training the Model

- **Alternative:** Directly train the model for a **supervised** task (e.g., node classification):



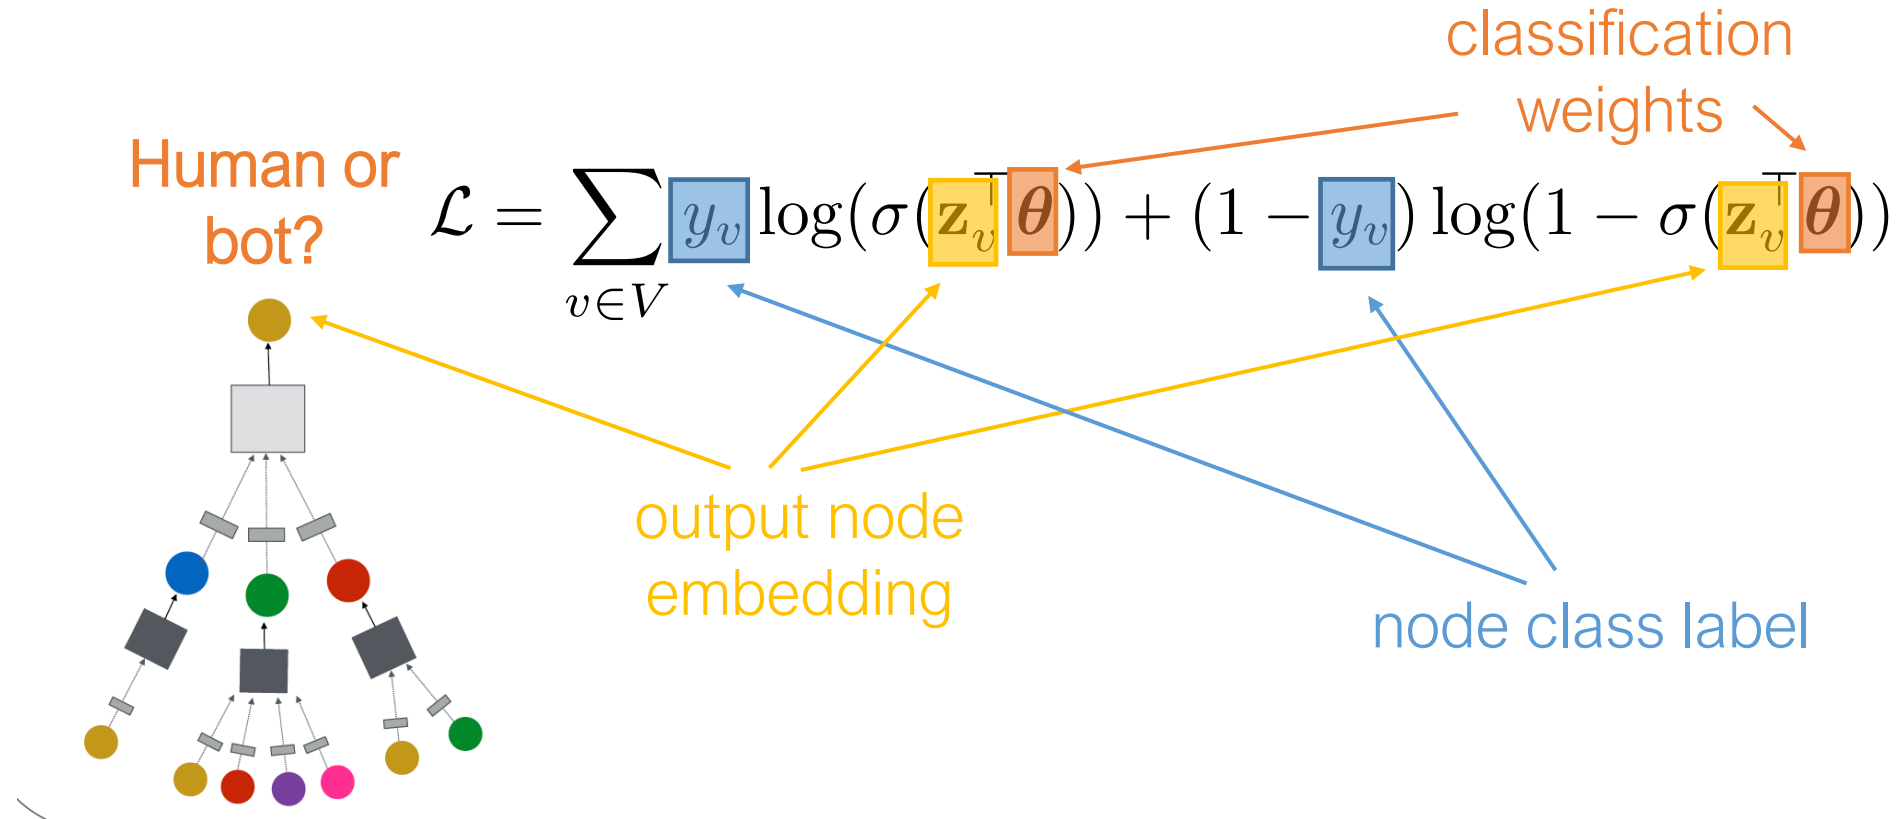
Human or bot?



e.g., an online social network


Training the Model

- **Alternative:** Directly train the model for a **supervised** task (e.g., node classification):




Overview of Model Design

1) Define a neighborhood aggregation function.





2) Define a loss function on the embeddings, $\mathcal{L}(z_u)$

Overview of Model Design




INPUT GRAPH

3) Train on a set of nodes, i.e., a batch of compute graphs



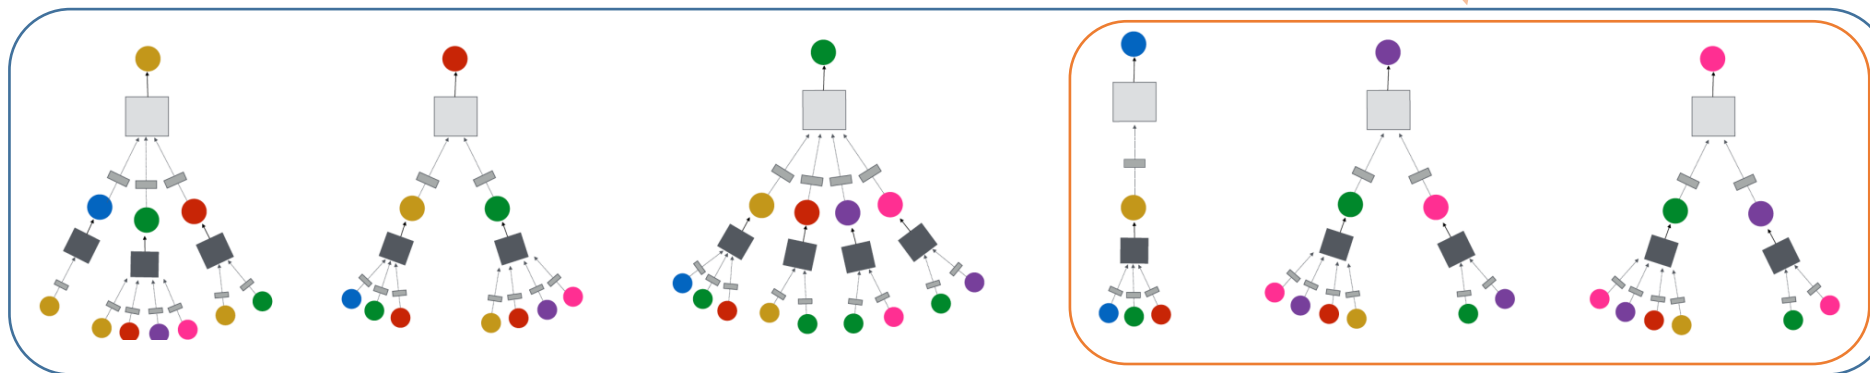
Overview of Model



INPUT GRAPH

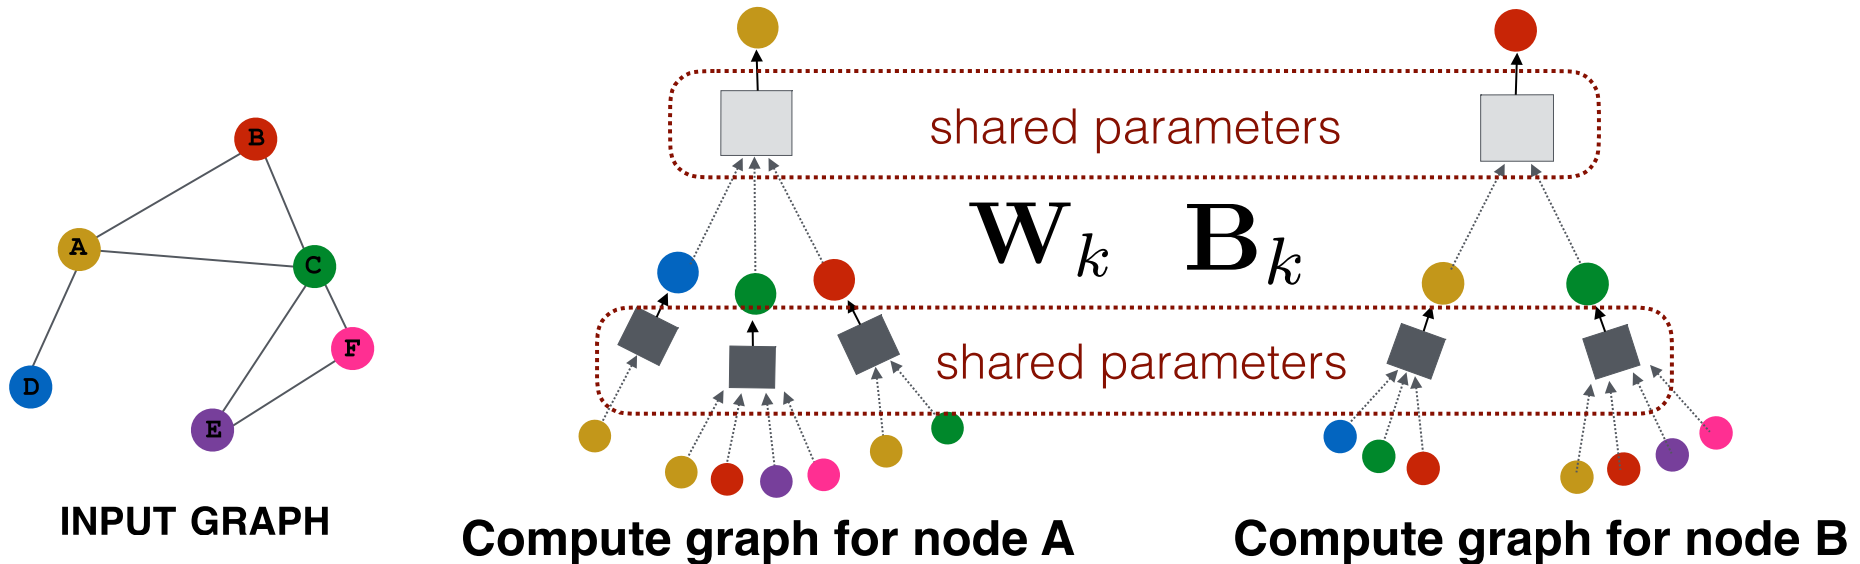
4) Generate embeddings for nodes as needed

Even for nodes we never trained on!!!!





Inductive Capability

- The same aggregation parameters are shared for all nodes.






Inductive Capability: New Graphs



Inductive node embedding → generalize to entirely unseen graphs

e.g., train on protein interaction graph from model organism A and generate embeddings on newly collected data about organism B

Inductive Capability



Many application settings constantly encounter previously unseen nodes.
e.g., Reddit, YouTube, GoogleScholar,


Need to generate new embeddings “on the fly”

Quick Recap


- **Recap:** Generate node embeddings by aggregating neighborhood information.
 - Allows for parameter sharing in the encoder.
 - Allows for inductive learning.

Neighborhood Aggregation

- Key distinctions are in how different approaches aggregate messages



What else can we put in the box?




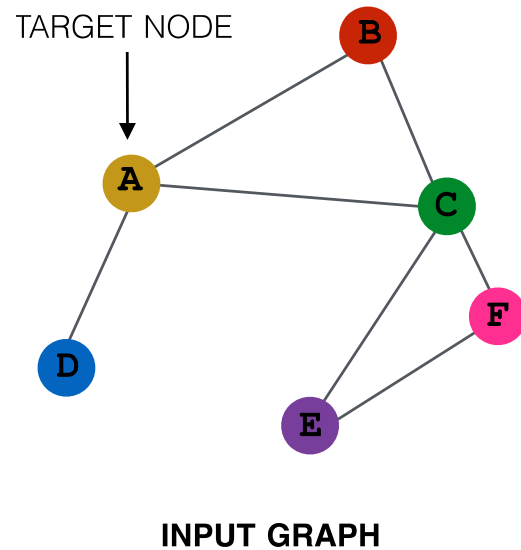
GraphSAGE

Based on material from:

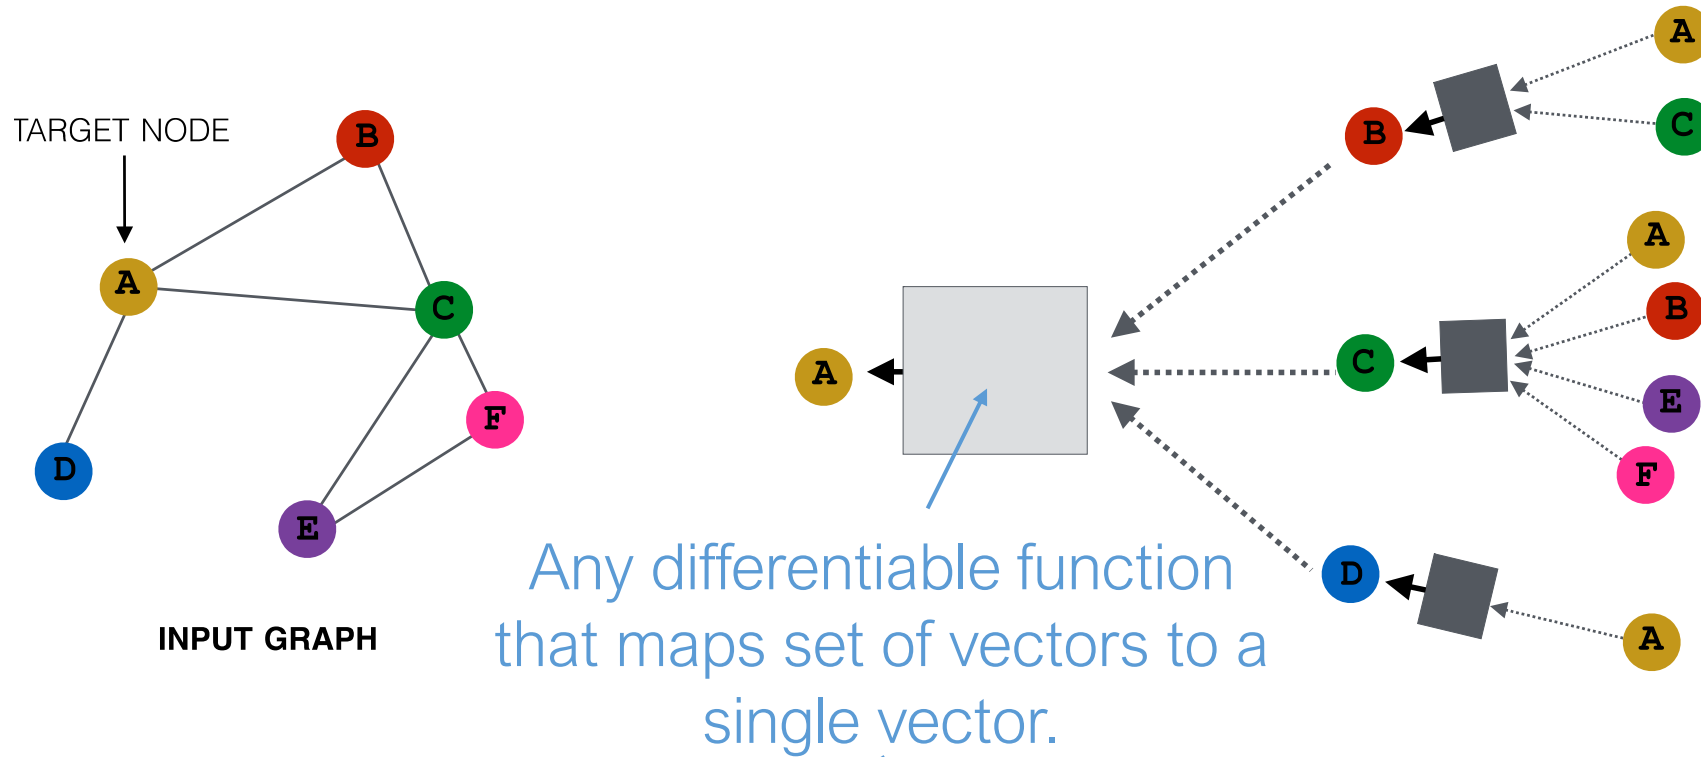
- Hamilton et al., 2017. [Inductive Representation Learning on Large Graphs](#). *NIPS*.

GraphSAGE Idea

- So far we have aggregated the neighbor messages by taking their (weighted) average, can we do better?



GraphSAGE Idea



$$\mathbf{h}_v^k = \sigma \left(\left[\mathbf{A}_k \cdot \text{AGG}(\{\mathbf{h}_u^{k-1}, \forall u \in N(v)\}), \mathbf{B}_k \mathbf{h}_v^{k-1} \right] \right)$$

Apply L2 normalization for each node embedding at every layer

GraphSAGE Differences

- Simple neighborhood aggregation:

$$\mathbf{h}_v^k = \sigma \left(\mathbf{W}_k \sum_{u \in N(v)} \frac{\mathbf{h}_u^{k-1}}{|N(v)|} + \mathbf{B}_k \mathbf{h}_v^{k-1} \right)$$

- GraphSAGE:

concatenate self embedding and
neighbor embedding

$$\mathbf{h}_v^k = \sigma \left(\left[\mathbf{W}_k \cdot \text{AGG} \left(\{ \mathbf{h}_u^{k-1}, \forall u \in N(v) \} \right), \mathbf{B}_k \mathbf{h}_v^{k-1} \right] \right)$$

generalized aggregation

GraphSAGE Variants

- **Mean:**

$$\text{AGG} = \sum_{u \in N(v)} \frac{\mathbf{h}_u^{k-1}}{|N(v)|}$$

- **Pool**

- Transform neighbor vectors and apply symmetric vector function.

$$\text{AGG} = \gamma(\{\mathbf{Q}\mathbf{h}_u^{k-1}, \forall u \in N(v)\})$$

element-wise mean/max

- **LSTM:**

- Apply LSTM to random permutation of neighbors.

$$\text{AGG} = \text{LSTM}([\mathbf{h}_u^{k-1}, \forall u \in \pi(N(v))])$$

Summary

- **The Basics**
- **GraphSAGE**
 - Generalized neighborhood aggregation.
- **Still a very active research field**